Context-Held-Out Perturbation Response

Biology / ML | Implementation and interview guide

Purpose and intended demonstration

Predict synthetic expression changes from fictional intervention descriptors and cellular contexts. This is a computational biology learning project with synthetic or fictional inputs, not a clinical or experimentally validated system.

How the implementation works

Generate responses from shared linear descriptor effects plus context interactions. Fit a multi-output ridge surrogate using intervention descriptors, context coordinates and pairwise interaction features. No actual gene targets or intervention sequences are supplied.

Evaluation and evidence

Train on six contexts and test on two unseen contexts. Compare with a training-mean response and a simpler descriptor-only model. Report expression RMSE rather than claiming real perturbation efficacy.

Context-Held-Out Perturbation Response

Biology / ML | Implementation and interview guide

Setup and commands

Requires Python 3.10 or newer and NumPy. Create a virtual environment: python -m venv .venv. On Windows use .venv\Scripts\python.exe if activation is unavailable. Install with python -m pip install -r requirements.txt. Run python app.py, then open results/report.html in your browser. Run python -m unittest discover -v to execute tests. The default seed is 42; change it with --seed. Replace sample_input.json using --input path-to-file.json, preserving its schema. Algorithm-only projects generate synthetic arrays from the seed rather than consuming external files.

Architecture and file responsibilities

CLI arguments -> JSON input or seeded generator -> domain pipeline -> validation or evaluation -> JSON and HTML artifacts. Imported functions can be used in a notebook or a separate application. Model checkpoints preserve learned arrays where applicable; see app.py for parameter order. This package does not include a production web service. Each project is independent and includes its own helper code so the ZIP can be extracted and run alone. app.py holds the experiment and domain functions: features, generate, run. kit.py contains reporting, input loading and only the retrieval or generation helpers needed by this project. sample_input.json contains editable inputs. test_app.py checks domain behaviors; test_reporting.py checks metrics and HTML output boundaries. results/metrics.json records outputs and results/report.html provides a browser-readable report. Numeric projects also save arrays or learned parameters to results/artifacts.npz; inspect its keys with numpy.load. MODEL_CARD.md explains intended use and limits. .github/workflows/tests.yml runs tests and the fixture experiment. No API key is needed.

Context-Held-Out Perturbation Response

Biology / ML | Implementation and interview guide

Limitations and next improvements

The context coordinates are known synthetic descriptors. Extrapolation in real biology may be far harder; no wet-lab intervention is designed or recommended. Use a licensed benchmark with context and perturbation splits, include control-adjusted metrics and compare simple baselines before larger models.

Interview walkthrough

Present the problem and demonstrate one normal example, then one failure or boundary case. Explain the baseline or fixture mode, the split or validation design, and one limitation before discussing an extension. Be ready to trace data through the functions listed above. Explain why the chosen metric matches the task and why the available results cannot establish production performance. Add your own documented experiment and be precise about what you implemented or changed. These are AI-assisted educational projects; filenames and synthetic scores are not evidence of client work.

Research context and reading

These topics were selected from research directions reviewed on 22 October 2026. Recent activity suggests useful skills to develop; it is not a forecast of hiring, commercial success or scientific validity. The implementation is a small educational baseline and does not reproduce or claim the performance of any cited foundation model. Reading: Benchmarking algorithms for generalizable single-cell perturbation response prediction (2025/2026): https://www.nature.com/articles/s41592-025-02980-0;
Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines (2025): https://www.nature.com/articles/s41592-025-02772-6.

Verified build results

project

"Context-Held-Out Perturbation Response"

test rmse

0.20623332065743902

descriptor only rmse

0.9859080242170135

mean baseline rmse

1.2673752950519062

heldout contexts

[6, 7]

scope

"Fictional intervention descriptors, not actionable genetic designs."

See results/metrics.json for all observations and TEST_RESULTS.txt for the executed tests. All datasets and replay responses shipped here are synthetic. Results can vary across seeds and numerical library versions.